

JUAN CARLOS GURDIAN

AI Engineer | LLM and Agentic Systems | Full Stack

Tampa, FL · (682) 414-7100 · juangurdian2003@gmail.com · jcgurdian.io · github.com/juangurdian

SUMMARY

AI Engineer shipping production LLM and agentic systems end to end, from architecture and backend to frontend and client delivery. Designed and led an autonomous outreach agent platform, built 4 of 11 production agents on a multi-tenant platform serving 24 client organizations, and shipped an MCP server exposing those agents as tools in Claude. CTO and cofounder of an AI fitness app in closed beta. English and Spanish.

WORK EXPERIENCE

yorCMO

AI Engineer

Sep 2025 – Present

Remote

AI-native fractional CMO firm operating the CORE Growth Platform for 24 client organizations.

- **Linky, platform lead.** Designed and led development of an autonomous LinkedIn outreach platform built on LLM agents. Sources roughly 6,000 target companies and 12,000 prospects per client organization from a defined ICP, researches each prospect and its recent news before writing, and generates messages matched to the user's own voice; about 90% of drafts are approved without edit. Directed one engineer.
- Built Linky's autonomous enrichment pipelines, scanning multiple sources to resolve each prospect's email, phone, LinkedIn profile, activity signals, and region, then scoring fit before outreach. Shipped the agentic loop to production: daily invite automation, acceptance detection, and a three-message cadence on human-paced timing behind a human-in-the-loop approval step, with user edits feeding back as training signal.
- **CORE Growth Platform.** Built 4 of 11 production LLM agents on a multi-tenant platform serving 24 client organizations across 25+ integrations (LinkedIn via Unipile, WordPress, HubSpot, Google Analytics and Ads, Mailchimp, Klaviyo, Apify, Notion, Shopify, Microsoft Graph), plus core platform structure and backend maintenance.
- Built and maintain an MCP server exposing platform agents as callable tools inside Claude, along with the shared orchestration and tool-calling layer the agents run on.
- **Reddit intelligence agent.** Ingests selected subreddits on a set cadence, scores threads for relevance, and drafts replies grounded in retrieved company resources, plus analytics on audience trends and content gaps.
- **PanelCast.** Meeting-to-content pipeline that builds per-user brand and voice profiles from social and writing samples, then generates on-brand blog posts with imagery, per-attendee LinkedIn posts, and short-form video cuts.
- **Practice Growth System.** Internal lead on a five-agent build automating content and referral outreach from client meeting transcripts (Fathom REST and webhooks, Microsoft Graph sending, approval queue). Delivered the pilot on schedule and onboarded a second organization with no firm-specific code.
- Implemented guardrails (PII redaction, role-aware prompts), evaluation via prompt regression and A/B tests, and RBAC with audit logs. Own backend services across Modal, Vercel, Supabase, and AWS with Docker, GitHub Actions CI/CD, observability, and per-agent cost and latency controls.

Gynka

CTO and Cofounder

Jan 2025 – Present

Remote

AI fitness coaching app, currently in closed beta.

- Architected and built the multi-agent AI coach engine running on AWS (FastAPI, Pydantic AI, Supabase, ARQ/Redis, RevenueCat) as sole engineer across backend, mobile, and infrastructure.
- Built the full iOS and Android app in React Native, distributed through TestFlight and in closed beta with advisors and early users. The coach ingests biomarker data from connected wearables and adapts each workout toward the user's stated goal.

AtomChat

Machine Learning Intern

Jul 2024 – Oct 2024

Remote

- Built and deployed a customer-facing LLM agent for automotive dealerships (OpenAI API, LangChain, Pinecone, MySQL, Google Calendar) serving WhatsApp, Instagram, Messenger, and Facebook.
- Implemented Scrapy ingestion pipelines and tuned vector database queries to keep retrieval real-time.

PROJECTS

Vault AI, open-source local-first AI platform · github.com/juangurdian/Vault-AI

- Built an LLM-based smart router that classifies each query into one of seven task types and selects the optimal local model in under 500ms, with regex fallback, response caching, and automatic model upgrades when context overflows the selected window.
- Next.js 16 / React 19 / TypeScript frontend over a FastAPI backend with Ollama, ChromaDB RAG, SQLite, ComfyUI image generation, and self-hosted SearXNG search; one-command Docker Compose deploy.

SKILLS

AI and Agents: LLMs, agentic workflows, RAG, evals, guardrails, prompt engineering, tool and function calling, MCP (server and client), multi-agent orchestration, LangGraph, LangChain, Pydantic AI, Anthropic and OpenAI APIs, OpenRouter, Ollama, pgvector, Pinecone, ChromaDB

Languages: Python, TypeScript, JavaScript, Java, SQL, Bash

Backend: FastAPI, Node.js, Express, Spring Boot, REST APIs, webhooks, OAuth, RBAC, PostgreSQL, MySQL, Supabase, Redis, ARQ

Frontend: React, Next.js, React Native, Vue.js, TailwindCSS, responsive UI

Cloud and DevOps: AWS, Modal, Vercel, Docker, edge functions, GitHub Actions CI/CD, observability, cost optimization

EDUCATION AND CERTIFICATIONS

Texas Christian University, B.S. Computer Information Technology

May 2025

AWS Certified Cloud Practitioner · Full Stack Web Development with MERN Stack and GenAI